

Viability of Twisted and Python for Application Server Herd Implementations

Viet T. Nguyen
University of California, Los Angeles

Abstract

This paper provides a cursory look into the use of the Twisted framework, built on top of the Python language, to build application server herds with particular interest to language considerations such as type checking, memory management, and multithreading.

1 Introduction

The motivation for this research comes from a desire at our company to build a “new Wikimedia-style service designed for news.” Due to differing requirements we sought an alternative to Wikimedia’s LAMP platform (GNU/Linux, Apache, MySQL, and PHP). Our service intends to provide much more frequent updates, access via a variety of protocols, and a robustness towards increasing mobile traffic share.

The requirements outlined for the service result in a number of high level technical desires. To avoid application server bottlenecks we intend to distribute the application servers. With distributed servers we require rapid diffusion of unstructured data via a flooding algorithm. Beyond data diffusion there are a number of “-ility’s” we hope to achieve: maintainability, reliability, scalability, parallelizability. Of course from a cost perspective we would also like ease of implementation.

2 Twisted

One of the many options available to implement such a framework is Twisted. Twisted is an event-driven networking engine written in Python [2]. Twisted shows promise because of its event driven. It also contains means to spawn processes beyond using multiple threads, allowing us to overcome the Global Interpreter Lock which will be discussed later in Section 4.2.




Figure 1: Topology of the network used to test the prototype.

3 Prototype

To assist in evaluating Twisted and Python a simple prototype was rapidly developed. The prototype implements a single application run in multiple instances to mimic multiple servers each running the application. They communicate with each other via TCP. The servers as a whole comprise a service. A client expects to be able to send data to the service and then request something based on the data it provided the service. Since a client does not necessarily interact with the same server the servers must propagate the client’s data to the other servers in the service. The topology of prototype servers is given by Figure 1. More information about the design of the prototype are given by [1].

Handling client input was straightforward in thanks to the event driven structure of Twisted. Polling or thread handling is unnecessary as the `twisted.internet.reactor` handles much of it. Implementation simply requires the inheritance of prototypical classes representing protocols (e.g. connections),

a means of spawning protocols (factories), and then the definition of desired behavior encapsulated in a method whose arguments are the incoming data from the client. From there Python provides numerous means for parsing data, including built-in string and list manipulation. Here the use of Twisted shines in quickly getting development to the actual desired functionality with little boiler plate code. Python's maturity and flexibility also prove advantageous here with its numerous capabilities via built-in structures and standard packages.

3.1 Simple Flooding Algorithm

Part of the prototype functionality requires that servers communicate in an intra-service manner to keep client data consistent across servers. To achieve this a simple flooding algorithm was implemented. Upon receiving of new client data the receiving server updates its internal store of client data (in memory) before forwarding the data to all of its known peers. When a peer receives this data it then updates its internal store of client data before proceeding to forward the data to all of its known peers sans the peer it received the data from. To prevent infinite recursive forwarding a peer checks the time stamp of the data and proceeds only if the time stamp is newer than the data in its local store.

Implementation is relatively straightforward except for having to understand the multi-layered callbacks required to establish a temporary connection to send data from server to server. Receiving and parsing of data from another server follows the same steps as receiving and parsing client data. In communicating with another server, however, an end point needs to be instantiated (e.g. `TCP4ClientEndpoint`), a connection made via the end point, and a call back provided to the connection. The connection then executes the callback providing the callback with the protocol that can then be used to actually send the data. As a result the method that performs data communication is separated from the location of desired communication by two object instantiations and a callback that must be executed by the reactor. When implemented the process works fine, but maintainability and readability suffers. However, in an event driven framework this problem is common.

3.2 Post-Mortem

The Twisted framework and Python very much succeeded in allowing us to quickly develop the prototype. It is our feeling that this setup can be used to develop our product in a scalable fashion.

Access via a variety of protocols is readily available due to the object-oriented design of the framework and its built-in implementation of numerous other protocols that can, more or less, be swapped with the instantiations used in the prototype. It appears to require little additional work to have servers communicate over `ssh` or `UDP`. The same applies for client communication. The design of the server herd addresses the concern for handling mobile traffic.

However, there are caveats to using Twisted and Python. An understanding of Python's multithreading implementation and memory management are important to avoiding development of software that scales and performs poorly. Some of these concerns are discussed in following sections.

3.3 Source Code

Source code for the prototype can be found at the following URL:

```
https://github.com/vietjtnghuyen/  
ucla-fall12-cs131-pr
```

Unfortunately the prototype lives up to its name. The prototype has not been designed for robust execution or even functionality. Instead it simply proves the ability to run multiple servers using Python and Twisted that can communicate with each other. It does not strictly adhere to good object-oriented design, don't-repeat-yourself principles (DRY), or other best practices.

4 Python

When discussing Python it is implicitly assumed that the discussion is in regards to CPython. This paper is no different, but an issue particular to CPython is relevant. Alternative implementations of Python exist and are discussed in Section 7. As such, we will explicitly refer to CPython here.

Part of our technical goals in selecting Twisted and CPython is parallelizability. Parallel software is generally executed using either multiple threads, multiple processes, or both. Threads can be considered light-weight processes that operate within a process, possibly under the control of the operating system [3].

4.1 Processes Versus Threads

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur

eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

4.2 Global Interpreter Lock

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

5 Twisted Revisited

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

5.1 Object-Oriented Design

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

5.2 Reactor

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

5.3 Processes

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue

diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

6 Alternative Event Driven Python Frameworks

Beyond the Twisted framework there are a number of other event based frameworks available for Python. These include `gevent` and `Eventlet`.

Both `gevent` and `Eventlet` operate based on green threads (greenlets in the case of `gevent`). These green threads are psuedo-threads in that they are threads implemented on the virtual machine rather than the operating system itself. Although these threads are lightweight, they appear to offer only a locking abstraction on top of regular threads. Regardless, both are based on threads rather than processes meaning they would not take full advantage of a multi-core server for parallel processing due to the GIL as previously discussed.

7 Alternative Python Implementations

Aside from alternatives to Twisted there are actually alternative implementations of Python. The standard implementation of Python is known as CPython. CPython, as its name suggests, is the Python language implemented in C as a bytecode interpreter. To address shortcomings of this implementation there are a number of alternative Python implementations such as Jython, PyPy, Stackless, and IronPython. Some of these alternatives are discussed below, including the shortcomings they intend to address.

7.1 Jython

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

7.2 IronPython

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue

diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

7.3 Stackless

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

7.4 PyPy

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Proin facilisis elementum sagittis. Praesent mattis augue in nulla ornare vitae dignissim nibh pharetra. Ut augue diam, mollis id porta vitae, convallis quis nisl. Curabitur eu nunc ac est faucibus elementum. Maecenas venenatis dictum iaculis. Phasellus ultrices, metus eget mattis adipiscing, est lorem molestie magna, eu bibendum sem turpis in mauris.

8 Blah Below

Sometimes you want to really call attention to a piece of text. You can center it in the column like this:

_1008e614_Vector.p

and people will really notice it.

The noindent at the start of this paragraph makes it clear that it's a continuation of the preceding text, not a new para in its own right.

Now this is an ingenious way to get a forced space. Real * and double * are equivalent.

Now here is another way to call attention to a line of code, but instead of centering it, we noindent and bold it.

size_t : fread ptr size nobj stream

And here we have made an indented para like a definition tag (dt) in HTML. You don't need a surrounding list macro pair.

fread reads from stream into the array ptr at most nobj objects of size size. fread returns the number of objects read.

This concludes the definitions tag.

References

- [1] EGGERT, P. Project. twisted twitter proxy herd. <http://cs.ucla.edu/classes/fall112/cs131/hw/pr.html>, November 2012.
- [2] LABS, T. M. Twisted. <http://twistedmatrix.com/trac/>, November 2012.
- [3] WIKIPEDIA. Thread (computing). [http://en.wikipedia.org/wiki/Thread_\(computing\)](http://en.wikipedia.org/wiki/Thread_(computing)), November 2012.